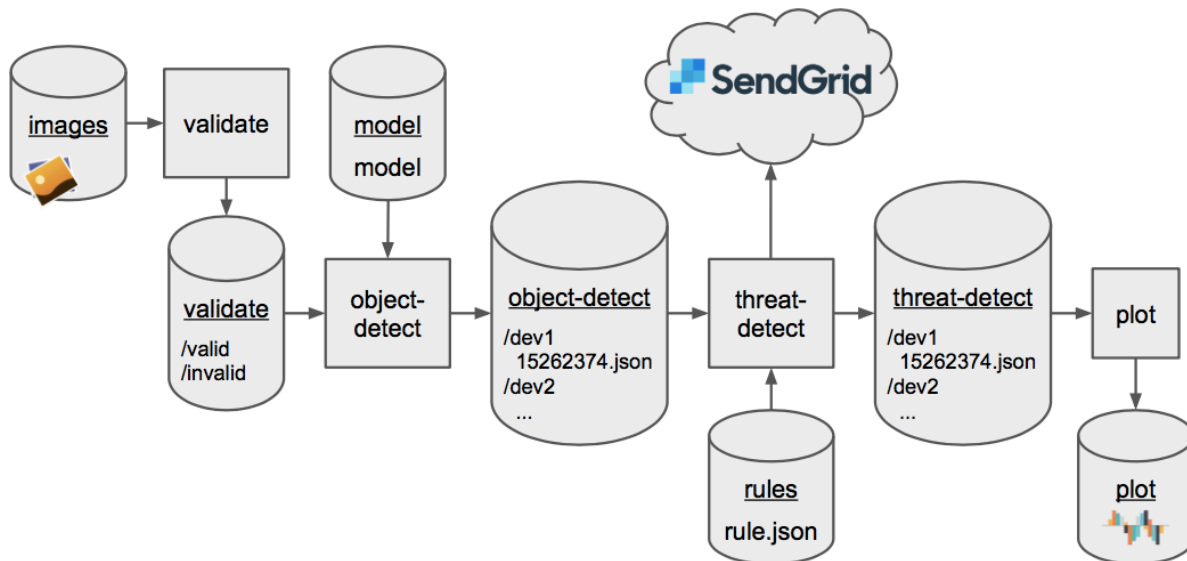
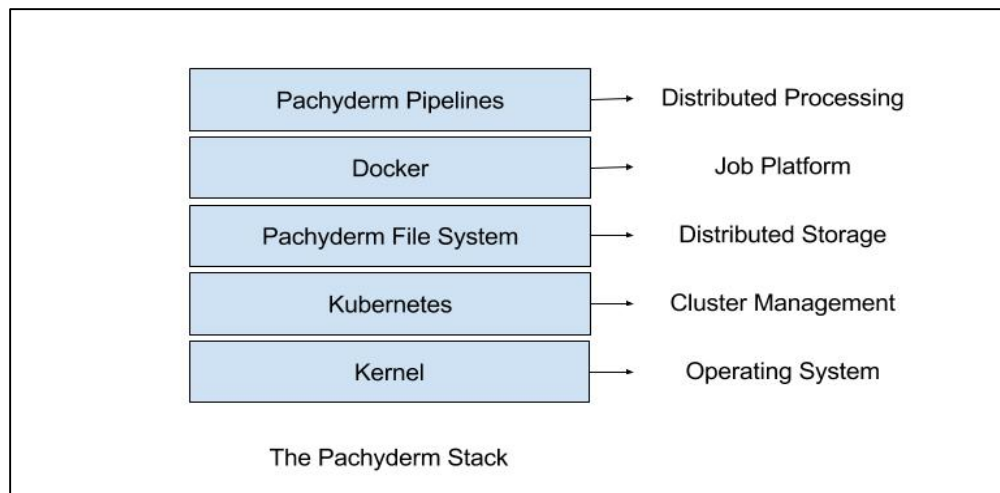


Putting it all together: Pipeline

So far, we had decided to implement a system which looks like this



The images for threat detection will be stored in a repository. The images will go through 'validate' process which will check images based on resolution and other parameters, and create file names. The 'validate' 'image' in the Docker container will check for validity of image. It will store the valid and invalid images separately. The TensorFlow model developed for object detection for different clients will be stored in 'model' repository. The Python 'image' will consist of codes that are executable using Python libraries in the 'TensorFlow' image in Docker container. The 'object detect' will process images and detect 'objects' in the image. The 'objects' detected in the image will be classified as threat or friendly based on rules in the file 'rule.json'. For the purpose of demonstration, we will consider only one rule. Different clients can have different threat perception and this will help determine if the 'object' in image is a threat or not. 'Threat-detect' 'image' in the Docker container will contain the package necessary to run for threat detection. The threat will be described in form of a json file with the date and time the threat was detected along with its location. This information will go to the user using an email service such as SendGrid. The threat once detected will be stored in the separate repository for regulatory purpose. This storage will happen in a 'static' 'object' storage platform like S3. For internal customer analytics, we will plot the number of threats against user or location. All the computation will happen on EC2 instance. Kubernetes will manage the EC2 instances (VMs). We will create a Pachyderm 'image' stored in Docker container called as 'pachd', that will automate the entire process of running different 'images'. It will under the hood manage Kubernetes and optimize the utilization of EC2 instance.



The image shows how we incorporate open-source software, working on cloud, managing and scheduling the processes using Pachyderm pipelines.

Implementation:

First, we login to a virtual machine which has Pachyderm, TensorFlow and Kubernetes installed on it. I executed `ssh pachrat@XXX.XXX.XXX.XXX` on my Mac Terminal. Once we login, we run the following commands.

kubectl get all – command that shows all that is running on Kubernetes.

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/dash	1	1	1	1	1h
deploy/etcd	1	1	1	1	1h
deploy/pachd	1	1	1	1	1h

NAME	DESIRED	CURRENT	READY	AGE
rs/dash-6c9dc97d9c	1	1	1	1h
rs/etcd-7dbb489f44	1	1	1	1h
rs/pachd-75c6c9b79f	1	1	1	1h

NAME	READY	STATUS	RESTARTS	AGE
po/dash-6c9dc97d9c-bz4kk	2/2	Running	0	1h
po/etcd-7dbb489f44-625d5	1/1	Running	0	1h
po/pachd-75c6c9b79f-rvqfq	1/1	Running	0	1h

NAME	IP	PORT(S)	TYPE	CLUSTER-IP	EXTERNAL-IP
svc/dash		NodePort	10.105.88.60	<none>	8080:30080/TCP,8081:30081/TCP
svc/etcd		NodePort	10.105.9.158	<none>	2379:32379/TCP
svc/kubernetes		ClusterIP	10.96.0.1	<none>	443/TCP
svc/pachd		NodePort	10.106.62.174	<none>	650:30650/TCP,651:30651/TCP,652:30652/TCP,999:30999/TCP

We see that 3 images are running. 'dash' will create a dashboard. 'etcd' tracks what data has been processed or any more data needs to be processed etc.

kubectl get pods - command that shows the status of pods

NAME	READY	STATUS	RESTARTS	AGE
dash-6c9dc97d9c-bz4kk	2/2	Running	0	1h
etcd-7dbb489f44-625d5	1/1	Running	0	1h
pachd-75c6c9b79f-rvqfq	1/1	Running	0	1h

pachctl version - command that shows the version of Pachyderm installed

COMPONENT	VERSION
pachctl	1.6.5
pachd	1.6.5

git clone https://github.com/dwhitena/mgmt690-pipeline.git- command that clones folder named mgmt690-pipeline from github.

```
Cloning into 'mgmt690-pipeline'...
remote: Counting objects: 71, done.
remote: Compressing objects: 100% (51/51), done.
remote: Total 71 (delta 21), reused 58 (delta 14), pack-reused 0
Unpacking objects: 100% (71/71), done.
Checking connectivity... done.
```

Checking the mgmt690-pipeline folder shows the images in that folder

```
pachrat@mgmt6:~/mgmt690-pipeline$ ls
object-detect  plot  README.md  threat-detect  validate
object-detect.json  plot.json  test_images  threat-detect.json  validate.json
```

This shows the object-detect, threat-detect, plot and validate folder with the json files in it. The test images folder will contain the images on which we will perform threat detection.

pachctl create-repo rules

pachctl list-repo

NAME	CREATED	SIZE
rules	5 seconds ago	0B
model	12 seconds ago	0B
images	About a minute ago	0B

The command above will create a repository for rules. List the repo to display the 3 repository.

wget

http://download.tensorflow.org/models/object_detection/ssd_mobilenet_v1_coco_11_06_2017.tar.gz

--2017-11-28 19:18:04-

- http://download.tensorflow.org/models/object_detection/ssd_mobilenet_v1_coco_11_06_2017.tar.gz

Resolving download.tensorflow.org (download.tensorflow.org)... 2607:f8b0:4006:813::2010, 172.217.12.144

Connecting to download.tensorflow.org (download.tensorflow.org)|2607:f8b0:4006:813::2010|:80... connected.

```
HTTP request sent, awaiting response... 200 OK
Length: 128048406 (122M) [application/x-tar]
Saving to: 'ssd_mobilenet_v1_coco_11_06_2017.tar.gz'
```

```
ssd_mobilenet_v1_co 100%[=====] 122.12M  326MB/s   in 0.4s
```

```
2017-11-28 19:18:05 (326 MB/s) - 'ssd_mobilenet_v1_coco_11_06_2017.tar.gz' saved
[128048406/128048406]
```

tar -xvf ssd_mobilenet_v1_coco_11_06_2017.tar.gz

```
ssd_mobilenet_v1_coco_11_06_2017/
ssd_mobilenet_v1_coco_11_06_2017/model.ckpt.index
ssd_mobilenet_v1_coco_11_06_2017/model.ckpt.meta
ssd_mobilenet_v1_coco_11_06_2017/frozen_inference_graph.pb
ssd_mobilenet_v1_coco_11_06_2017/model.ckpt.data-00000-of-00001
ssd_mobilenet_v1_coco_11_06_2017/graph.pbtxt
```

The commands above downloads TensorFlow and relevant libraries.

pachctl put-file model master -c -f frozen_inference_graph.pb – Command puts model file in model repository with the graph.

pachctl list-repo – Command that lists the Pachyderm repositories

NAME	CREATED	SIZE
model	3 minutes ago	27.83MiB
rules	3 minutes ago	0B
images	4 minutes ago	0B

To modify the rule.json file, we will first navigate to that folder and then add the email address to which alert emails are to be sent. The following commands are used to do that.

pico rule.json

cat rule.json

```
{
  "from_email": "support@baimsecurity.com",
  "to_email": "pgujela@purdue.edu",
  "threat_classes": [1, 3, 4, 8]
}
```

cat validate.json – Command to check the contents of validate file

```
{
  "pipeline": {
    "name": "validate"
  },
  "transform": {
    "image": "dwhitena/mgmt-validate",
    "cmd": [ "/bin/ash" ],
    "stdin": [
      "python /code/validate.py $images /pfs/out/valid /pfs/out/invalid"
    ]
  },
  "parallelism_spec": {
    "constant": "1"
  }
}
```

```

},
"input": {
  "atom": {
    "repo": "images",
    "glob": "/*"
  }
}
}
}

```

We can see from the codes that valid and invalid images are stored in different folders.

pachctl create-pipeline -f validate.json

pachctl create-pipeline -f object-detect.json

pachctl create-pipeline -f threat-detect.json

All the above commands will execute the respective files in Kubernetes by Pachyderm. First, we will validate the image, detect objects and then detect threats.

kubectl get pods – This command will display all the Docker images that Kubernetes is running

NAME	READY	STATUS	RESTARTS	AGE
dash-6c9dc97d9c-bz4kk	2/2	Running	0	1h
etcd-7dbb489f44-625d5	1/1	Running	0	1h
pachd-75c6c9b79f-rvqfq	1/1	Running	0	1h
pipeline-object-detect-v1-9hgwm	2/2	Running	0	2m
pipeline-threat-detect-v1-d68kf	2/2	Running	0	41s
pipeline-validate-v1-z7fb9	2/2	Running	0	4m

pachctl list-job – This command will list out the files that Pachyderm is running and its status

ID	COMMIT	PROGRESS	DL	UL	STATE	STARTED	DURATION	OUTPUT	RESTART
a768938b-33fc-4d2d-922c-d2aeb111dcb1	plot/48c11ccad0c64d348e85e3e18bcae9b6	4 minutes ago	Less than a second	0	1 + 0 / 1	102B	13.67KiB	success	
57bd23eb-2484-4d5a-a2cf-b47cd20c3151	threat-detect/74fec69fa6cf472f9684c2dcb8aef916	7 minutes ago	Less than a second	0	1 + 0 / 1	219B	102B	success	
4d2706d3-1fb2-44d6-af11-f5392efbe72f	object-detect/11ae39f89a4c4d9ca5abd4cf365002f1	7 minutes ago	8 seconds	0	1 + 0 / 1	27.86MiB	102B	success	
ac9bbddc-38a1-4a5f-a004-921a6d87cc3f	validate/7cf96f8d97994a75a7faeeb2b8b3589e	7 minutes ago	Less than a second	0	1 + 0 / 1	31.11KiB	31.11KiB	success	

We can see that all files have ran successfully. We have even finished executing the plot command. We can see the files and repositories created in Pachyderm along with the graph.

PACH DASH

Home

Recent Changes

Repos

Pipelines

Jobs

Settings

7 repos

3 Manually Ingested • 4 Computed Outputs

plot

Computed Output Repo

Updated 4 minutes ago

1 files • 0 dirs • 13997 B • 1 commits

threat-detect

Computed Output Repo

Updated 6 minutes ago

0 files • 1 dirs • 102 B • 1 commits

object-detect

Computed Output Repo

Updated 6 minutes ago

0 files • 1 dirs • 102 B • 1 commits

validate

Computed Output Repo

Updated 6 minutes ago

0 files • 1 dirs • 31852 B • 1 commits

images

Manually Ingested Repo

Updated 6 minutes ago

1 files • 0 dirs • 31852 B • 1 commits

rules

Manually Ingested Repo

Updated 18 minutes ago

1 files • 0 dirs • 117 B • 1 commits

model

Manually Ingested Repo

PACH DASH

Home

Recent Changes

Repos

Pipelines

Jobs

Settings

Search Pachyderm

X

4 pipelines

0 starting • 4 active • 0 restarting • 0 paused

plot

Active Pipeline

Updated 4 minutes ago

0 active jobs • 1 inputs • 1 output commits

threat-detect

Active Pipeline

Updated 10 minutes ago

0 active jobs • 2 inputs • 1 output commits

object-detect

Active Pipeline

Updated 12 minutes ago

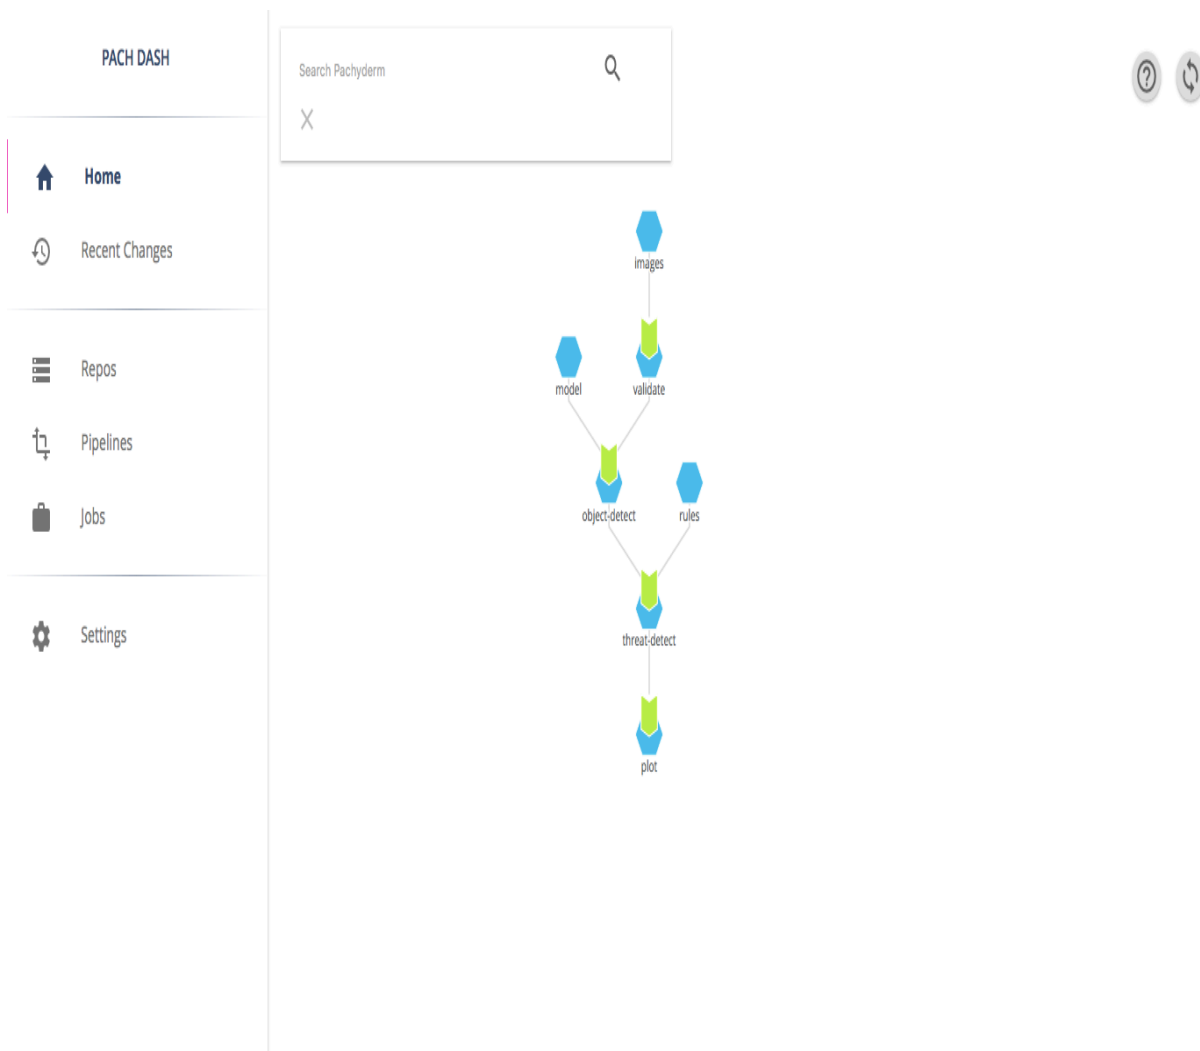
0 active jobs • 2 inputs • 1 output commits

validate

Active Pipeline

Updated 14 minutes ago

0 active jobs • 1 inputs • 1 output commits



[SUSPECT SPAM]Alert! Possible Threat Detected by BAIM Security



support@baimsecurity.com <support@baimsecurity.com>

Gujela, Prasad Vasant

Tuesday, November 28, 2017 at 4:48 PM

[Show Details](#)

Check your security system. A possible threat has been detected!

The image will show the alert mail received after processing the threat image from the company 'BAIM Security'. In future, we can add image that was used to detect threat in the email. A better way of alerting would be through text or call on cell phone. We so far tested the architecture on single image. Next time, we will productionize it to run many images automatically.